

Process_Management.md

Process Management

A process is an instance (i.e., running) of a program. Or simply a program in execution is called as a process. While a package/application/program is an executable file lying in some directory in the hard drive (storage), a process is started when the user or another program calls (initializes) the program and this is opened in RAM and CPU allocates time to this.

There are fundamentally two types of processes in Linux:

- **Foreground processes:**
Also known as interactive processes. Foreground processes are initialized and controlled through a terminal session.
- **Background processes:**
Also known as non-interactive/automatic processes. Background processes are not connected to a terminal; they don't expect any user input.

Every process has a parent. The top most process is `init`, whose process id is 1

Stages of a process:

1. User-running
2. Kernel-running
3. Ready to run in memory
4. Asleep in Memory
5. Ready to run, swapped
6. Sleep, Swapped
7. Pre-empted
8. Created
9. Zombie.

ps

`ps` is the command to check all the processes and their resource utilization (CPU & memory)

```
kk@kmaschine:~$ ps
  PID TTY          TIME CMD
 79004 pts/1        00:00:05 zsh
 79008 pts/1        00:00:00 zsh
 79038 pts/1        00:00:00 zsh
 79040 pts/1        00:00:00 zsh
 79041 pts/1        00:00:00 gitstatusd-linu
 79583 pts/1        00:00:00 bash
 79589 pts/1        00:00:00 ps
```

`ps aux` gives processes started by all the users on the Computer

```
USER      PID %CPU %MEM    VSZ   RSS TTY      STAT START   TIME COMMAND
root         1  0.0  0.0 312232 11164 ?        Ss   Nov30    0:24 /sbin/init splash
root         2  0.0  0.0      0      0 ?        S    Nov30    0:00 [kthreadd]
root         3  0.0  0.0      0      0 ?        I<   Nov30    0:00 [rcu_gp]
root         4  0.0  0.0      0      0 ?        I<   Nov30    0:00 [rcu_par_gp]
root         6  0.0  0.0      0      0 ?        I<   Nov30    0:00 [kworker/0:0H-events_highpri]
root         9  0.0  0.0      0      0 ?        I<   Nov30    0:00 [mm_percpu_wq]
root        10  0.0  0.0      0      0 ?        S    Nov30    0:00 [rcu_tasks_rude_]
root        11  0.0  0.0      0      0 ?        S    Nov30    0:00 [rcu_tasks_trace]
root        12  0.0  0.0      0      0 ?        S    Nov30    0:03 [ksoftirqd/0]
root        13  0.0  0.0      0      0 ?        I    Nov30    1:34 [rcu_sched]
root        14  0.0  0.0      0      0 ?        S    Nov30    0:00 [migration/0]
root        15  0.0  0.0      0      0 ?        S    Nov30    0:00 [idle_inject/0]
root        16  0.0  0.0      0      0 ?        S    Nov30    0:00 [cpuhp/0]
```

PID - Process ID

TTY - Controlling Terminal

STAT: Process State Code

TIME: Total time of CPU Usage

CMD: The command that triggered the process

RSS: Both swap memory and physical Memroy

VSZ: Virtual memory usage of the process

%CPU: CPU Time used by the process run time

%MEM: Processes set size to the physical memory on the machine

START: Process Start time

Various Process States:

D - Uninterruptible sleep

I - Idle kernel thread

R - Running or runnable (in the queue)

S - Interruptible sleep (waiting for event or input)

T - Stopped by job control signal

t - Stopped by debugger

X - dead (generally, not seen)

Z - Zombie

Kill

`kill` is the command to stop any process

`kill <pid>`

To force kill a program :

`kill -9 <pid>`

Top

This command gives real time information about the processes and their utilization

The data gets refreshed every 3 seconds. You could change it by pressing 's'

The processes are sorted by CPU utilization by default. If you want to change that you could press 'f' and select other field.

```
top - 00:03:24 up 7 days,  4:39,  1 user,  load average: 1.48, 1.25, 1.29
Tasks: 354 total,   1 running, 353 sleeping,   0 stopped,   0 zombie
%Cpu(s):  9.3 us,   3.4 sy,   0.0 ni, 87.2 id,   0.0 wa,   0.0 hi,   0.1 si,   0.0 st
MiB Mem : 15848.2 total,  3396.5 free,  4716.1 used,  7735.6 buff/cache
MiB Swap: 31470.5 total, 31470.2 free,    0.2 used.  9016.0 avail Mem
```

PID	USER	PR	NI	VIRT	RES	SHR	S	%CPU	%MEM	TIME+	COMMAND
75648	kiran	20	0	3404676	796924	158544	S	16.2	4.9	23:10.44	Web Content
77611	kiran	20	0	29.7g	267280	102352	S	12.9	1.6	5:41.45	spotify
69361	kiran	20	0	5703412	430796	111456	S	12.3	2.7	5:37.43	gnome-shell
69946	kiran	20	0	991540	163572	104052	S	10.6	1.0	0:07.16	flameshot
77538	kiran	20	0	3783640	235952	146576	S	9.3	1.5	3:17.51	spotify
68686	kiran	20	0	694304	113252	69568	S	8.3	0.7	5:40.31	Xorg
68587	kiran	20	0	2878944	24364	18220	S	7.6	0.2	3:58.56	pulseaudio
75399	kiran	20	0	36.9g	354600	207564	S	7.6	2.2	8:04.43	signal-desktop
72831	kiran	20	0	5697188	904224	366876	S	4.0	5.6	30:33.72	GeckoMain
77560	kiran	20	0	1132156	122484	83636	S	3.6	0.8	1:55.59	spotify
75370	kiran	20	0	497740	126128	82396	S	3.3	0.8	2:57.41	signal-desktop
73075	kiran	20	0	2918256	216436	131400	S	2.0	1.3	2:15.20	Web Content
69533	kiran	20	0	325380	11836	6984	S	1.7	0.1	0:22.18	ibus-daemon
79828	root	20	0	13120	5724	5116	S	1.7	0.0	0:00.05	systemd-hostnam
78123	kiran	20	0	589160	76200	45696	S	1.3	0.5	0:15.65	terminator
801	root	20	0	636612	7872	6668	S	0.7	0.0	1:53.29	systemd-logind
72994	kiran	20	0	2780836	300808	110468	S	0.7	1.9	3:27.43	Web Content
79810	kk	20	0	22868	4240	3260	R	0.7	0.0	0:00.21	top
1	root	20	0	312232	11164	7828	S	0.3	0.1	1701:22	systemd

Shift + p -- Sorts the processes by highest CPU utilization

Shift + m -- Sorts the processes by highest Memory utilization

Fields in the Header

us: Amount of time the CPU spends executing processes for people in 'user space'

sy: Amount of time spent running system's kernel space's processes.

ni: Amount of time spent executing processes with a manually set nice value.
id: Amount of CPU idle time.
wa: Amount of time the CPU spends waiting for I/O to complete.
hi: Amount of time spent servicing hardware interrupts.
si: Amount of time spent servicing software interrupts.
st: Amount of time lost due to running virtual machines ('steal time')

Fields in main output

PID: Shows task's unique process id.
USER: User name of owner of task
PR: Stands for priority of the task.
NI: Represents a Nice Value of task. A Negative nice value implies higher priority, and positive Nice value means lower priority
USER: User name of owner of task.
SHR: Represents the amount of shared memory used by a task.
VIRT: Total virtual memory used by the task.
%CPU: Represents the CPU usage.
TIME+: CPU Time, the same as TIME, but reflecting more granularity through hundredths of a second.
%MEM: Shows the Memory usage of task.

Sorting

By default top sorts the entries by CPU usage.
Press **M** (upper case) to sort by Memory
To revert to sort by CPU usage press **P** (upper case)

Priority (PR) & Niceness (NI)

Priority and Niceness are related. While Priority is

PR is the priority level. The lower the PR, the higher the priority of the process will be.

NI is niceness of the process. The higher the niceness, the process will get lower precedence.

PR value can be computed by the following formula: $PR = 20 + NI$.

the process with niceness 3 has the priority 23 ($20 + 3$) and the process with niceness -7 has the priority 13 ($20 - 7$). You can check the first by running command `nice -n 3 top`. It will show that top process has NI 3 and PR 23. But for running `nice -n -7 top` in most Linux systems you need to have root privileges because actually the lower PR value is the higher actual priority is. Thus the process with PR 13 has higher priority than processes with standard priority PR 20.

Niceness ranges -20 to 20. Lesser the Niceness higher the priority.

`nice` and `renice` commands are used to update Priority of the user processes. PR of only user processes can be altered with `nice` and `renice` commands

nice: A program can be started with a specific niceness with the below command

```
nice -n <nice_value> ./myProgram
```

Example

```
nice -n 9 ./myscript.sh
```

renice: Priority of a program can be altered with renice command

```
renice -n <nice_value> <PID>
```

Example

```
renice -n 9 1344
```

load average :

Load Average is the measure of the load on the processors. Load Average has three fields each of which is the number of processes waiting for CPU in the last 1 minute, last 5 minutes, and last 15 minutes.

For a Computer with 1 CPU, the load average should always be less than 1 and above 0.9 indicates high utilization of CPU. For a computer with 2 CPUs, the uptime should be under 2 and so on so forth

Further Reading

[Understanding ps command](#)

[A guide to Linux TOP command](#)